Attack-resilient media using phone-to-phone networking

P.W.G. Brussee

Attack-resilient media using phone-to-phone networking

by

P.W.G. Brussee

to obtain the degree of Master of Science
at the Delft University of Technology,
to be defended publicly on Wednesday November 30, 2016 at 10:30 AM.

TUDelft

Contents

1

Introduction

Modern media and news distribution is shifting from traditional media to social media. Modern media production and information distribution is shifting from traditional outlets. Also media consumption is shifting from traditional outlets to digital and mobile devices. From con-sumer to pro-sumer. News is not bound to expert opinion or an editorial news desk. The opinion of the individual can be broad-casted without the need for a professional studio and equipment. No office is required anymore to link a mobile eye-witness to a mass audience.

Given this context mobile devices are increasingly important and smart-phone in particular. A smart-phone has the unique property of being a ubiquitous device that is highly mobile and extremely connectible. Most smart-phones even have one or more cameras to produce multi-media content that can be shared immediately from the device. Finally the entire world has a smart-phone. Even or especially areas without tradition infrastructure they are ubiquitous.

In crisis situations, like natural disaster or unrest, the smart-phone becomes particularly important, exactly for the earlier mentioned properties, as under such conditions utilizing centralized infrastructure is undesired (censorship) or physically impossible. Decentralized can work in these situations. Censorship and large scale monitoring is difficult in decentralized networks.

Tribler is a fully decentralized video-on-demand system. [10, 13, 15] It is autonomous, attack-resilient and self-organizing. [1, 12] It uses network overlays called communities to offer features like keyword search and managing contributions to channels for discover-ability of content. It offers privacy through layered encrypted tunnels similar to the TOR network.[3, 4, 17]

So far Tribler only supports desktop and server versions of Linux, Mac and Windows. The necessity of moving to mobile devices calls for Tribler functionality to be enabled to run on these resource limited devices. In this thesis the first prototype is presented that has all Tribler functionality fully enabled on mobile devices. The prototype is build for Android OS since Android dominates the smart-phone market.

Sharing opinion and discussion is enabled on social media. A global dialog is possible through social media like Twitter, Instagram and Facebook. People receive local and global news perceived relevant to their group on their social media feed. Young people in particular are shifting to social media as their number one source. [9] As shown in this age distribution graph, figure 1.1.

Figure 1.1: Main news sources split by age [9]

Not only is social media more and more becoming a major distribution channel for news, it also starts delivering input for news and story creation. This way social media has become both the source and outlet for investigative journalism.

A smart-phone has the unique property of being a ubiquitous device that is highly mobile and extremely connectible. Figure 1.2 shows that world-wide 1,4 billion smart-phones were sold to end-users last year. And this number keeps rising. Even or especially areas without traditional infrastructure they are ubiquitous.

Figure 1.2: Number of smart-phones sold to end users worldwide from 2007 to 2015 (in million units) [14]

The capabilities and versatility of smart-phones enable it to be used for production and consumption of news and social media. The users themselves are turning from con-sumers into pro-sumers [?]. With regards to production most smart-phones have one or more cameras to record multi-media content that can be shared immediately from the device. Eye-witnesses often have smart-phones at hand to immediately record an event with and post it on social media [?]. No news desk or professional equipment is necessary to relay news directly from eye-witnesses to the masses anymore [?]. People with camera phones can reach millions of people with multi-media in a very short timespan, becoming social journalists. The ease of reaching a global audience by an individual with a smart-phone diminishes the role of an expert curator handling incoming information.

Also with regard to the consumption medium mobile devices increasingly replace the role of traditional outlets like TV, newspapers and other physical media.

In crisis situations, like natural disaster or unrest, people need to communicate and coordinate their efforts to restore safety. In this context the smart-phone becomes particularly important because it is often carried on person and provides connectability

In recent calamities [**?** **?**] people could mark themselves as safe on social media, effectively broadcasting that information to all their family and friends on social media instead of contacting them one by one or not at all due to congestion in the communication channels.

However, several natural disasters have taken out the necessary infrastructure on numerous occasions for a prolonged period of time.

Therefore we require a solution to enable social media on smart-phones that does not require infrastructure.

The main research question is: how feasible is it to run all Tribler functionality on mobile devices?

Secondary question is: using the unique properties of mobile devices, what features can be added or enhanced? Secondary question is: considering the constraints and opportunities of mobile devices, what unique ability can be utilized to extend or enhance Tribler?

<div style="text-align: right;">

2

</div>

Problem description

2.1. Privacy and censorship

Pervasive monitoring of digital citizens by Internet providers on behalf of governments to enforce censorship laws raises severe privacy concerns. The lack of anonymity becomes a problem when the users privacy is being invaded. Revealing personal information can be deduced from search queries for example, or associations on social platforms. When this information can be used for targeted advertising it becomes very valuable, and creates an incentive for the parties that have access to this information to sell it to third parties. Social media companies use targeted advertisement as part of their business model. Information considered private by users of social media is actually used to broker targeted advertisements. Subsequently users can be confronted with their information being misused in various ways beyond their control. This lack of control over your own privacy can lead to arbitrary interference as defined in UDHR article 12. Integration of social media on regular websites makes every page-view and click on these websites traceable to an individual, directly benefiting the business model of targeted advertisements

In fact the business model of social media appears to be serving targeted advertisements to its users on behalf of third parties. What's even worse is social media integrated into regular websites to de-anonymize and track the whereabouts of users even outside of the social media realm. Whenever users lose control over their privacy it becomes a serious problem.

Figure 2.1: Viral spreading from one device to another within an off-line region

The incentive to de-anonymize the user, not only causes a lack of privacy, but also a potential lack of freedom of expression, as it hands key information to a censor: who is expressing dissent and who is associated

<div style="text-align: center;">5</div>

with this person on-line. Cyber-suppression has become a reality when you no longer can be associated with opinion-makers or foreign journalists on-line.

Internet exchange (IX) infrastructures are among the central components in the inter-network architecture that are also vulnerable to monitoring, censorship and Internet kill-switches. As such, not everyone has unrestricted access to the Internet due to censorship and surveillance. In fact a significant part of today's Internet users is affected by these attempts to hide or distort reality. This interference directly affects the universal right to freedom of opinion and expression as stated in article 19 of the Universal Declaration of Human Rights (UDHR).

For example: Arab Spring. Kill switches are real. [**?**]

What are kill-switches?

Large portions of the global dialog on social media is uncontrolled by traditional media or governments.

Utilizing infrastructure is undesired because of censorship.

The sophistication of censorship techniques is pushed forward by the drive to stay ahead of attempts trying to circumvent it. Increasingly though, Internet traffic is put under surveillance and obfuscation techniques are targeted by restrictions.

2.2. Adversary model

From the Arab Spring scenario we know Internet kill switches are real, so we must assume the existence of a powerful adversary. The following threats [11] have been identified for similar circumstances:

- The adversary can observe, block, delay, replay, and modify traffic on the underlying network. Thus end-to-end security must not rely on the security of the underlying network.

- The adversary has a limited ability to compromise smart-phones or other participating devices. If a device is compromised, the adversary can access any information held in the device's volatile memory or persistent storage.

- The adversary can choose the data written to the transport layer by higher protocol layers.

- The adversary cannot break standard cryptographic primitives, such as block ciphers and message-authentication codes.

We assume the adversary cannot eavesdrop, jam, delay, replay, modify or spoof wireless communication between smart-phones. The adversary cannot compromise smart-phones or other participating devices.

2.3. Distributed solutions

To ensure that no controlling party can exercise censorship authority must be distributed over all users. If all information is located in one or a few places, the parties in charge of that location will still have control over it, so information must distributed all users as well, creating a *communication* system. Finally, all users must be able to share, order and appreciate information of other users, in other words the essence of social media: social interaction. With everyone being able to interact in the same way we need to distribute functionality over all users, creating a *cooperation* system. Fully distributed systems capture the characteristics just mentioned. To render the effect mute of the Internet kill switches in existence shown to be used we propose a distributed solution. Without any central component in the system it is no longer susceptible to censorship without everyone participating. Leveraging the properties of mobile devices, like smart-phones as stated above, together with the features of Tribler we see a perfect match. Because censorship and large scale monitoring is difficult in decentralized networks our proposed solution can work in these situations.

Networks that are not fully decentralized can be disrupted by taking down a limited number of nodes, less than the total number of nodes. Thanks to a server-less design we can say: The only way to take Tribler down is to take the entire Internet down.

Figure 2.2: Various partial solutions to mend the broken Internet according to www.youbroketheinternet.org

Mobile devices are essential to take the next step in the face of censorship. If we can use smart-phones to create a fully decentralized social network we can say: The only way to take Tribler down is to take everybody's smart-phones away.

The importance of mobile devices in this context is crucial , if we want to do even better in spreading information, which is not been done before, we surmount than next hurdle.

Peer-to-peer communication technology is essential for a server-less distributed system. Mobile devices typically do not require infrastructure to exchange information, like those equipped with Bluetooth or capable of ad hoc Wi-Fi. Smart phones are ubiquitous everywhere in the world and used to access social media and retrieve information from the Internet. Fortuitously these are also the type of mobile devices that can communicate peer-to-peer.

Example: Arab spring [**?**]

Various initiatives have been started to deal with one or both of these problems. [**?**]

Figure 2.2 shows a mapping of projects that are or have been working on that. The fragmentation is clear from the figure. There are a few big players, as can be seen from the figure, but none of them provide a full solution.

There is no de-facto solution available on mobile platforms. [**?**] Previous research has shown that there is no solution that solves the problem entirely in a sustainable way. Tribler is our attempt to solve this problem

entirely in a sustainable way. To see if it is feasible to apply it in a mobile context as described above, we need to port it.

2.4. Contributions

The main contribution of this thesis is making Tribler available on mobile devices and the utilization of their unique properties in the context of censorship and communication during crises. This work also enables a new direction for future research with Tribler: mobile devices. The resilience of Tribler against attacks on the Internet is greatly increased. The goal of Tribler is to become an information sharing platform that protects the privacy of its users and is resilient to attacks while not relying on existing infrastructure.

Previous attempts failed to deliver all functionality. [**?** **?** **?**] Maintainability issues with earlier designs were a large part of the reasons why. We changed the architecture of Tribler for our approach.

The scientific contributions of this work are the experiments to verify the feasibility of Tribler on mobile devices and the ability to perform research with Tribler fully geared towards mobile devices

3

Tribler functionality

Tribler is a fully decentralized video-on-demand system. [10, 13, 15] It is autonomous, attack-resilient and self-organizing. [1, 12] It uses network overlays called communities to offer features like keyword search and managing contributions to channels for discover-ability of content. It offers privacy through layered encrypted tunnels similar to the TOR network.[3, 4, 17]

3.1. Video-on-demand

Tribler introduces a server-less video-sharing platform with privacy enhancing technologies that provides a Youtube-like social media experience. Video-on-demand means that users can simply click and play videos in a streaming fashion, so without waiting for the entire video to be present on the device. You can search from within the application and browse for videos in channels rated automatically by popularity. Simply clicking on a video and watching it while streaming is supported via the BitTorrent-aware Tribler video server and integrated video player VLC.

3.2. Self-organizing

The BitTorrent protocol is used to download and upload the content. A BitTorrent swarm can use the "distributed sloppy hash table" (DHT) of a specific torrent to discover peers and gather meta-information. Also part of the DHT protocol is the self-organizing behavior of maintaining a routing table of known good nodes. Tribler uses these features to coordinate the exchange of videos and meta-data fully automatically. Users do not have to manage any files or configuration manually at all to be active on the platform.

3.3. Autonomous operation

New content discovery can be discovered automatically via a "Channel"-community that users can subscribe to. Communities are network overlays used by Tribler to offer functionality like search, add, remove and comment on content. To discover all channels in existence Tribler is subscribed to the so called "AllChannel"-community and "Search"-community by default. These are used to exchange information about what channels and torrents are out there and who likes them and knows about them. Each channel has its own community that rules permissions and meta-data of content. These communities operate autonomously and are transparent to the user, who only sees channels and search results show up in the GUI.

3.4. Attack-resilience

The server-less technique of Tribler is resistant to large scale monitoring and censorship, because there is no central point that can be controlled to gather or block information easily. To monitor or censor the network effectively on a large scale you need control of a significant number of the communication links. Censorship does not have an effect if the majority of users does not cooperate with the censor. The beauty of a fully distributed design is that communicating directly between peers, or via a local network, works without the need for external communication links. This is why the server-less technique of Tribler is resistant to Internet kill-switches as well because, even if the attacker can block all communication-links, users can always connect

off-grid. Such kill-switches are typically deployed for the purpose of censorship, but won't stop a connectible device, like a laptop, from physically moving. No network infrastructure is required for viral spreading of the entire video platform.

These properties will ensure social media with resilience against Internet kill switches, natural disasters and censorship.

3.5. Anonymity

Tribler can protect the privacy of users by hiding their identity. To connect to others on the network anonymous connections are created on behalf of downloaders and uploaders. By routing the network traffic over a circuit of multiple hops it becomes difficult to trace the origin and destination. This way the privacy of users remains protected while they actively participate on the platform

Multiple encrypted tunnels are layered such that every consecutive tunnel from the initiator to a relay is going through the previous tunnel. Every relay works on behalf of its predecessor so no relay knows the identity of the initiator save for the first relay. Since all communication is properly encrypted no relay can perform a successful man-in-the-middle attack.

This is similar to how Tor works, except Tribler uses UDP rather than TCP for performance reasons. The capability of hiding your identity is greatly advantageous to the user if his or her human rights are violated, like free speech.

However, this privacy feature requires a lot more bandwidth of the network than without anonymity: a ratio of 13 GB for every anonymous 1 GB of data. Since bandwidth is limited and transitory it can be beneficial to exchange unused bandwidth for a promise of bandwidth in the future, or another reward. The research group behind Tribler is currently building a fully decentralized accounting system and open exchange market using block-chain technology with the purpose of building trust on-line and creating the Internet of Money.

3.6. Towards Tribler on mobile devices

In the previous sections we discussed the features and applications of Tribler. So far Tribler only supports desktop and server versions of Linux, Mac and Windows. The necessity of moving to mobile devices calls for Tribler functionality to be enabled to run on these resource limited devices. All positive and negative properties that come from these features are transferred to mobile devices, which will add their own distinct properties to the mix.

3.6.1. Opportunities

What does mobile allow in addition to desktop? Mobile devices are inherently easy to move around and very portable. In case of a breakdown in communication infrastructure the mobile devices with wireless radio transmitters can still connect ad hoc and moved within range if necessary. Via WiFi a device can connect to the Tribler network via existing infrastructure or other peers via ad-hoc WiFi. Using NFC Tribler can start a Bluetooth connection and transfer the installation package peer-to-peer. Also via NFC Tribler can exchange channel ID's to subscribe to another Tribler channel peer-to-peer. The same method can be used to exchange bootstrap peers for the Tribler network. The peer-to-peer WiFi Direct file transfer would be much faster than Bluetooth.

To bring Tribler to mobile devices will give potentially millions of users access to these features on the move. Expanding the Tribler network with mobile devices would also benefit the research that can be performed on the live network.

3.6.2. Challenges

What does working on mobile phone mean? The portability of mobile devices requires any network interface and power supply to be wireless. Mobile devices are typically equipped with batteries to operate without a power cord. Considering the size and weight the capacity is limited. Smart-phone batteries usually barely hold a charge that would sustain a day of heavy usage. Tribler could potentially drain the battery much faster than normal. Heavy encrypted network traffic not only demand constant radio transmissions, but also CPU processing. In case of hidden seeding building circuits of 3 tunnels with layered encryption quadruples the amount of crypto work. Because Multichain punishes cheating like double spending by a permanent ban, it must never loose information and flush everything to permanent storage before continuing. Mobile devices typically have flash memory with limited write-cycles compared to classic hard drives that are commonly

found in desktop computers.

4

Design and architecture

We now present a design to address the challenges and utilize the opportunities of working on a mobile device as put forward in the previous chapter. First in terms of functionality and other requirements, second the overall system architecture.

Previous attempts at designing a mobile version of Tribler failed to properly separate re-usable components. This resulted in unmaintainable code and increased difficulty in testing. We use a top down approach to come to a high level system architecture and in determining the reusable components or the current Tribler application. Some functionality of Tribler has been shown to work on Android before [? ? ?]. Our design will fill in the gap of connecting the pieces of the grand puzzle to make Tribler fully work on mobile.

4.1. Functional requirements

Following from the problem description and Tribler functionality described in Chapter 3, we define the following functional requirements:

A1. The implementation must be capable of effortless peer-to-peer transfer of itself.

A2. The implementation must be capable of publishing videos to other devices without the need for an Internet connection.

A3. Any video available on the device must be directly publishable.

A4. The implementation must enable a user to record videos.

A5. The implementation must enable a user to create a channel.

A6. The implementation must enable the owner of an existing channel to edit it.

A7. The implementation must enable creating a torrent file from a file available on the device.

A8. The implementation must support streaming video playback.

A9. The implementation must support all other Tribler functionality not mentioned above.

4.2. Non-functional requirements

Following from the challenges and opportunities, as described in Section 3.6, we define the following non-functional requirements:

B1. The mobile device must be capable of running Tribler independently.

B2. The implementation must incorporate the existing Tribler Python core and the required C/C++ libraries.

B3. The implementation must consider the restricted resources of a mobile platform in terms of RAM and processing power.

B4. The implementation must support WiFi, Bluetooth and NFC peer-to-peer features.

B5. The implementation must utilize the built-in camera for recording videos.

B6. The mobile device must be connectible via WiFi or mobile data connection

B7. Search results ought to appear in less than 1 second for local content.

B8. The implementation must be distributed as a single installable container.

B9. The implementation must be able to keep running in the background even if the user is not actively using it.

B10. All processing tasks must be performed asynchronously.

B11. The user must be able to interact with the implementation via a graphical user interface (GUI).

B12. All ongoing tasks must be indicated as such in the GUI.

B13. The GUI must stay responsive to the users' input while performing a background task.

B14. The GUI must stay responsive to the users' input while presenting large amounts of data on screen.

B15. If invalid input is provided by the user through the GUI the user must be asked to correct the input.

B16. If a recoverable error occurs, the implementation must automatically retry.

B17. If an exception occurs, the user must be able to restart Tribler.

B18. Upon restarting, the implementation must return to a working state.

B19. The entire build tool-chain must be integrated on the build server of Tribler.

B20. As much code as possible must be covered by tests.

B21. The implementation must be agnostic to version differences of supported platforms and operating systems.

B22. The user interface design must follow established best practices.

B23. The implementation must be attack-resilient.

4.3. System architecture

The requirements dictate how the system architecture of Tribler on mobile devices will take shape. The proposed architecture as shown in Figure 4.1 is specifically designed for maintainability and clearly separates components with a distinct responsibility.

Figure 4.1: System architecture design

The requirement (B2) of re-using the Tribler Python core and its C/C++ dependencies requires the architecture to incorporate these as is. Mobile platforms are not required to support Python code natively and the implementation must be distributed as a single installable package and run independently (req. B1). As a consequence, a Python interpreter must be incorporated into the implementation. The user must interact with the implementation via a GUI, as stated by requirement B11. As such a GUI is not part of the Tribler core, it has to be included separately in the design. A separate GUI can be made and optimized for any specific platform and target device. For instance, a design can be made for large surface displays and another one for small touch-screens, like a smart-phone. This leads to the design choice of creating an API, that allows for a common interface across all platforms, to let the GUI communicate with the existing Tribler core. The API will yield a more maintainable solution than previous attempts to bring Tribler to mobile devices [? ? ?] that did not include such an API.

5

Implementation

In Chapter 4 we proposed a generic design to enable Tribler on mobile devices. Many types of connectible mobile devices exist. In the context of the problem description from Chapter 2, devices used for human communication, such as smart-phones, will be our prime focus. The largest potential user base in the smart-phone market can be reached by targeting Android OS [5]. Therefore, our implementation will target Android OS.

Following the choice for Android, the generic design in Figure 4.1 can be further specified as presented in Figure 5.1. Two main additions are a separate video player and a Java native interface (JNI) component. In the following sections, each component from Figure 5.1 will be described in more detail. The complete build tool-chain is presented in Section 5.5, which combines everything to build the final Android application package (APK). Finally, the statistics of the implementation are presented in Section 5.6.
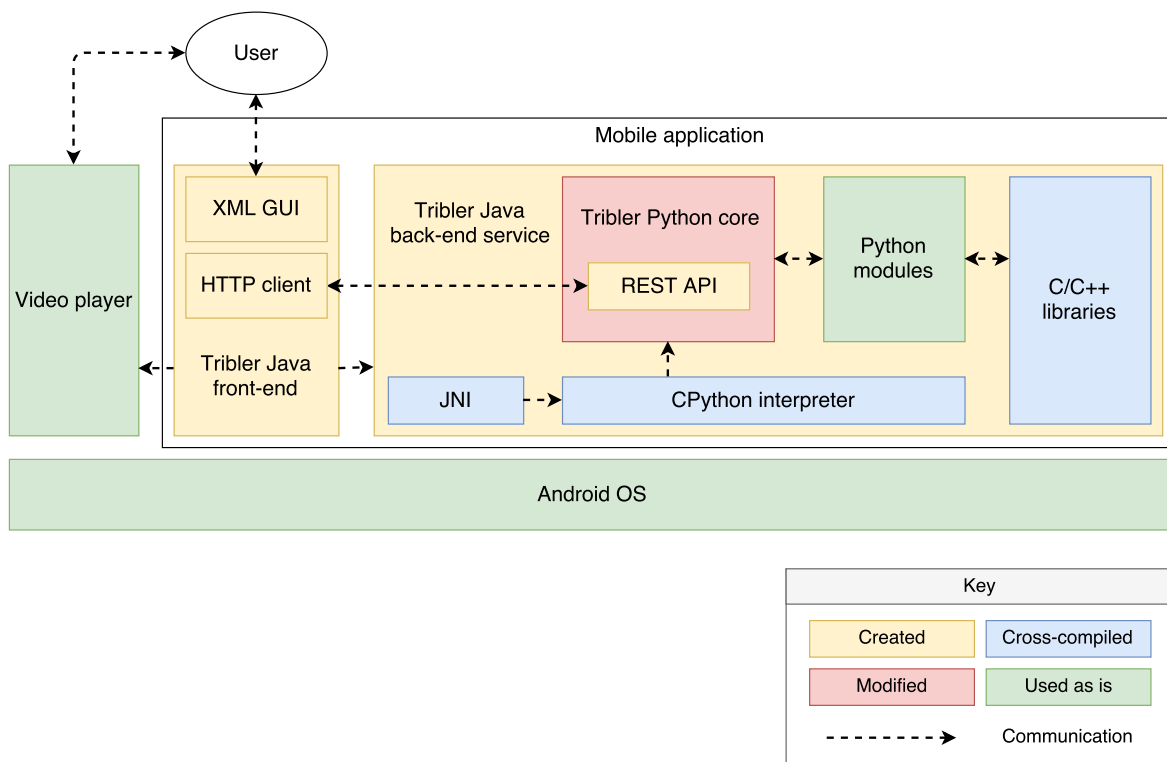
Figure 5.1: Implemented system architecture

5.1. Android OS

Android is an operating system, based on the Linux kernel, that runs on smart-phones, tables, wearables and smart-TVs. It provides a Java Virtual Machine (VM) and a Java application framework API.

Since API version 14, near field communication (NFC) and Wi-Fi peer-to-peer (P2P) connections between compatible devices is supported. And since API version 16, a NFC push message can be used to start large file transfers over Bluetooth (req. A1). This can be used in the context of the problem description, for Tribler needs to be able to spread wireless from phone to phone. Using this NFC push message functionality, the transfer of the APK file is fully automated (req. B4). Figure 5.4 shows the instructions for the user of this effortless transfer: just hold the phones back to back. Bluetooth, rather than Wi-Fi, is the technology of choice here, because the former has a standard file transfer protocol, built into Android OS, and the latter does not. This means there are no prerequisites on the receiving device for receiving and installing Tribler from a nearby phone via NFC+Bluetooth. If NFC is not available, all other options to transfer the APK are presented to the user to choose from instead. If NFC is available, but not enabled on the device, the implementation will prompt the user to do so. In that case, the other options are accessible via a button with the Bluetooth icon on the action bar. NFC is used to enable easy sharing of channels as well, for example your own channel or your favorites. After receiving the NFC push message, Tribler is automatically started and asks the user if the received channel must be added to their favorites.

Our implementation supports API version 18 and higher, because of reasons explained in Section 5.2.5. 85.6% of Android devices run API version 18 or higher [7]. Android support libraries are used to abstract from differences between API versions (req. B21).

Inter-application communication, via Android intents, must to be secured (req. B23) . An Intent is a messaging object to request an action from another app component [8]. Therefore, all Android intents are explicit for internal actions and the action of all received intents is checked, especially broadcasted intents [2]. Also, only the activities and services that should be publicly accessible are exported in the application manifest.

5.2. Tribler Java back-end service

Our implementation uses Java to build upon the Android Java API. To run an application in the background, even when the screen is turned off, it must be run as a service. Therefore, every component, except the GUI, is part of the back-end service (req. B9). The service is started as a separate process by the Java front-end (req. B13). This way the user can be presented with the GUI, that indicates the service is loading in the background (req. B11, B12).

5.2.1. Tribler Python core

Tribler is written entirely in Python, but Android does not support this natively. Python is an interpreted language, Because our design requirements (B1, B2) we incorporate a Python interpreter into our design. On top of that the entire core of Tribler can run, containing a REST HTTP API module also written in Python.

All communication with the front-end is done asynchronously via a REST API (req. B10).

Finally, the REST API communicates with an HTTP Client on the user interface side via JSON.

5.2.2. JNI

Using Java requires the use of the Java native interface (JNI) to communicate with the C/C++ components. JNI enables functions that are written in C to be callable from Java and vice-versa. The Python interpreter is started by the Java service using JNI, after loading the necessary C/C++ libraries.

5.2.3. CPython interpreter

Tribler uses C++ implementation of torrent protocol, called libtorrent. To use this C/C++ library from within Python a certain capable interpreter is required. P4A offers such an interpreter: CPython.

What alternatives are there, besides P4A, to run python code on Android? 0. QPython: scripting, cannot build regular .apk 1. QtAndroid, inmiddels alternative, niet gereed toen wij hieraan werkte (juni 2016, qt 5.7 android service) 2. PGS4A: no longer in development 3. SL4A: no longer in development Why is P4A chosen? Because this project continued from previous work (legacy code) build on P4A. Even though the revamp version was build from scratch, it still is the best choice because it not only provides the Python interpreter, it also is a complete tool-chain to cross-compile native libraries with bindings and build a standalone Android app installation package (.apk).

5.2.4. Python modules

Tribler is written entirely in Python, but most of its dependencies are written in C/C++. To use these libraries on a mobile device they need to be compiled for the right embedded-application binary interface (EABI) including all nonstandard dependencies. Figure 5.2 shows the dependency tree.
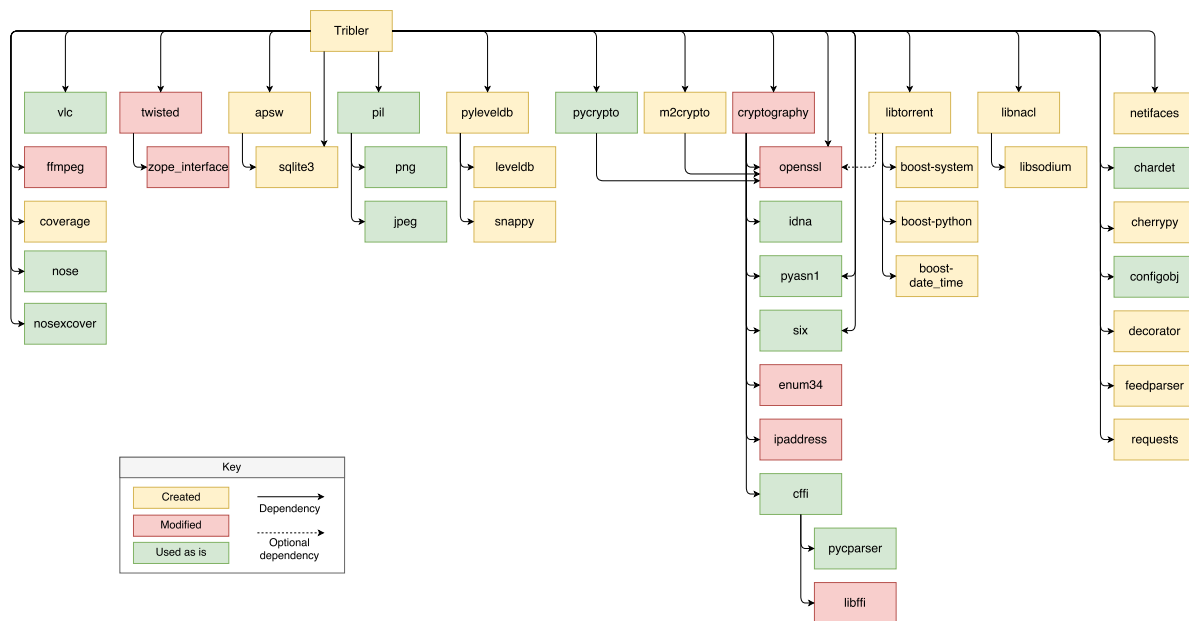

Figure 5.2: Tribler dependencies in terms of Python-for-Android recipes

5.2.5. C/C++ libraries

The standard C library of Android differs from the GNU C Library (glibc), which makes it not trivial to port Linux libraries to Android. All C/C++ dependencies of Tribler were therefore linked against glibc and glibc is included as a shared library (req. B1, B2). Static linking could result in unexpected behavior if more than one library is linked [6], and Tribler uses many, as shown in Figure 5.2. Since Android 4.3 (API version 18) shared libraries do not have to be loaded in order manually anymore. Therefore,

The Python code loads other Python and native modules directly.

Calling C code directly from Python is possible by using the Python ctypes module to load a native dynamic-link library (.so files on Android) or by using the Python/C API of CPython. This API enables a library to define functions that are written in C to be callable from Python. These Python bindings are the glue between pure Python and pure C code. SWIG can generate the boiler plate code for this. Libtorrent, one of Triblers' main components, uses Boost.Python to provide a standard C++ API on top of the Python/C API.

The Python/C API is actually so powerful it even provides access to the internals of the interpreter to mess with the global interpreter lock (GIL) which could be released during native C calls to improve the multi-threading performance of Tribler crypto.

5.3. Tribler Java front-end

The requirements on asynchronous communication (B10) and responsiveness (B13) require the decoupling of the GUI from the back-end.

The GUI is created by a native Android Java application, which talks to the REST API module. The API combined with Java is better for parallelization than the coarse grained locking by the CPython interpreter. Shared memory threading in Python code is restricted to single-thread performance because of the global interpreter lock (GIL).

Having two separate Python interpreters in distinct processes talking to each other, because that means using a very resource heavy Python GUI instead of the regular and lightweight native Android Java XML GUI. The latter has tools available for automated UI testing.

Figure 5.3 shows the main menu of the Tribler app. The first three items presented to the user are the main views: the user's favorite channels, their own channel and discovering popular content. The next four items

Figure 5.3: Navigation menu of the Tribler app

Figure 5.4: NFC+Bluetooth transfer of app or channel

Figure 5.5: Channel with newly discovered content

Figure 5.6: Search results showing video content

are actions related to creating new content or related to the application. It is organized this way to benefit from the simplicity of a flat menu while the actions are grouped by context with a header for clarity.

Users can browse through a list of popular channels or their own favorites. Each channel has an indicator of the amount of users that have added that channel to their list of favorites, as can be seen behind the menu in Figure 5.3. Channels contain multi-media content added by their respective channel owner or everyone, depending on the security policy of that channel. Tribler allows three settings: open, semi-open and closed.

The favorite channels and popular channels views both show a list of channels contain content

Search is accessible from any channel view. Upon typing keywords into the search bar debounce

5.3.1. HTTP Client

The HTTP client that talks in JSON with the REST API is build on the popular library OkHttp and fits perfectly to RxJava with a library called Retrofit. The Retrofit library enables a very declarative API client.

Because of the nature of Android to destroy interface elements if a configuration chance occurs, the asynchronous tasks running in the background must be registered and unregistered properly to avoid memory leaks. To detect notorious memory leaks on Android we use another library made to do exactly that: LeakCanary. Android provides a way to deal with continuity of activities with fragments that can remain in memory which were used in conjunction with composite subscriptions of RxJava.

5.3.2. XML GUI

Due to the fact that Android targets mobile devices it is very optimized for low resource usage. Therefore memory is freed more aggressively and the application is often paused or stopped and restarted if the user switches to another app. Running two interpreters with Python Kivy front-end and Python Tribler back-end would be doubling memory usage for the interpreter itself and require an inter-process protocol as well. Native Java with XML front-end and Python Tribler back-end brings the user experience seamlessly in line with the native UI.

To run in the background Tribler uses an Android service and all communication is performed asynchronously. The reactive programming paradigm is a perfect fit for asynchronous tasks. Thanks to RxJava and RxAndroid asychronous multi-threaded coding is made very enjoyable: As shown in the code example performing IO tasks on the dedicated Android thread and making UI changes on the main thread becomes trivial.

5.4. Video player

To support streaming playback of videos (req. A8) a capable video player is required. The video player VLC is integrated in the desktop version of Tribler. However, such a library, and integration of a custom GUI, is hard to maintain. VLC is offered as a standalone Android application package (APK) from their website [16]. Therefore, rather than implementing our own GUI, this APK is embedded as a whole inside Tribler's APK as an asset (req. B8). If VLC is not yet installed on the device, the user is prompted to do so and offered the version from inside Tribler. This allows the user to install VLC without further requirements.

5.5. Build tool-chain

The Python-for-Android tool-chain uses recipes to cross compile the C/C++ libraries with Python bindings with the necessary build tools. These recipes are like a high level make file.

Figure 5.7: High level overview of the build tool-chain

The libraries are compiled for any ARMv7 compatible platform, and little effort is required to replace Android with another OS in this respect.

5.6. Implementation statistics

All requirements have been met. The implementation consists of: 12.806 + X lines of code in the Tribler repository in 35 pull requests and Y lines of code in the Python-for-Android (P4A) repository in 46 pull requests and 2 lines of code in the M2Crypto repository in 1 pull request. This excludes the work of reviving the old P4A toolchain.

Of the X lines (x %) consists of unit testing code. Y lines (y %) is made up of Android specific libraries, and the remainder is.. Table 5.1X shows the 25 largest source file contributions for this thesis work.

Code coverage and testing details are further explained in chapter 6.

LOC	File	Path
718	MainActivity.java	.../org/tribler/android/MainActivity.java
666	MyChannelFragment.java	.../org/tribler/android/MyChannelFragment.java
482	MyUtils.java	.../org/tribler/android/MyUtils.java
403	DefaultInteractionListFragment.java	.../org/tribler/android/DefaultInteractionListFragment.java
318	start.c	android/TriblerApp/app/src/main/jni/src/start.c
314	TriblerViewAdapter.java	.../org/tribler/android/TriblerViewAdapter.java
281	build.gradle	android/TriblerApp/app/build.gradle
256	IRestApi.java	.../org/tribler/android/restapi/IRestApi.java
234	ChannelFragment.java	.../org/tribler/android/ChannelFragment.java
186	AssetExtract.java	.../org/kivy/android/AssetExtract.java
182	BeamActivity.java	.../org/tribler/android/BeamActivity.java
175	ListFragment.java	.../org/tribler/android/ListFragment.java
172	CopyFilesActivity.java	.../org/tribler/android/CopyFilesActivity.java
166	AndroidManifest.xml	android/TriblerApp/app/src/main/AndroidManifest.xml
166	EventStreamCallback.java	.../org/tribler/android/restapi/EventStreamCallback.java
155	ChannelActivity.java	.../org/tribler/android/ChannelActivity.java
150	ViewFragment.java	.../org/tribler/android/ViewFragment.java
149	PythonService.java	.../org/kivy/android/PythonService.java
149	SearchActivity.java	.../org/tribler/android/SearchActivity.java
141	EditChannelActivity.java	.../org/tribler/android/EditChannelActivity.java
131	FilterableRecyclerViewAdapter.java	.../org/tribler/android/FilterableRecyclerViewAdapter.java
130	BaseActivity.java	.../org/tribler/android/BaseActivity.java
114	SearchFragment.java	.../org/tribler/android/SearchFragment.java
112	TriblerDownload.java	.../org/tribler/android/restapi/json/TriblerDownload.java

Table 5.1: Top 25 of largest source file contributions.

6

Verification and validation

To verify if running all Tribler functionality on mobile devices is feasible we take a look at relevant performance characteristics. We take several measurements to quantify how well it works in terms of our non-functional requirements.

For scale a laptop and PC are included in some measurements. We focus on the design aspects and evaluate the key performance indicators:

- *usability* in terms of latency and startup time,

- *resource usage* in terms of CPU utilization,

- *performance* in terms of frequency and duration of function calls,

- *scalability* in terms of discovery time and bandwidth accounting.

6.1. Content discovery

New content is generated on the device, like for example a video that has been recorded. Before anyone can view that content it has to be added to a channel and discovered by other devices. We measure the amount of time it takes for other devices, which are subscribed to the channel, to discover new content. We also measure the amount of time it takes for content to be added to a channel in the form of a torrent. The amount of time it takes for a torrent to be created is relative to the size of the content. Therefore these measurements are normalized to a file of 1 MB. Depending on the random walk in the channel's community content can be discovered either very quickly or after a while due to property of eventual consistency. On one Nexus 6 device a channel is created to which 13 other devices subscribe via NFC. Then repeatedly a new video is recorded and added to that channel. A content discovered event is registered on each device individually and all devices are synced with NTP. Each device is connected to the same WLAN and within 1 to 2 meters distance from the access point.

Figure 6.1: All devices used in the experiment showing the about Android screen.


Figure 6.2: Sequence of events with time values from median of results.

Figure 6.3 shows the amount of time it takes a number of devices to discover new content on a subscribed channel.

Figure 6.3: Elapsed time from adding new content to a channel and it being discovered by subscribed peers.


Figure 6.4: Time required to create a torrent and add it to a channel, normalized for 1 MB of content.

The results show that within 4 seconds 9 devices have discovered the new content. From the 10th device and beyond an increase in discovery time is noticeable. This could be explained by the fact that only 10 peers are connected at the a time. Much more peers are needed to give an accurate representation in terms of scalability. What can be concluded from this experiment is that on the same local network the first device discovers the new content in less than 2 seconds. From the 10th device onwards the dissemination slows down.

6.2. Multichain performance

Multichain is the new accounting system of Tribler. This feature is central to the concept of trust in the Tribler network and is very important for the future as other functionality will be built upon it. If mobile devices are to become full-fletched nodes on the network they must support this feature. With Multichain any peer registers the bandwidth it exchanges with other peers. It aggregates these exchanges in blocks and signs them like a receipt and sends that to the other party to sign as well. These blocks are linked in a block-chain to foil attempts of cheating the system.

Multichain is explicitly designed to not use a global state. Syncing a global state is less scalable in terms of storage and network bandwidth.

The creation and signing of these blocks is measured to determine if it scales well. Multichain signs a block every 10 minutes, meaning our experiment of generating 25.000 blocks represent about half a year (173,6 days) of continuous effort. The database containing these blocks will grow over time, but should not slow down too much because of it. Measurements were taken on six different devices on multiple moments during development. Figure 6.5 show the performance graphs of every measurement.

Figure 6.5: Creating and signing of 25.000 blocks between two peers

Clearly visible from the graphs is that Multichain does not scale linearly on any device. They also show that mobile devices are at least a factor of two slower than an ordinary laptop and scale worse. Due to the nature of block-chain every new block needs to contain the hash value of the previous block. If a database lookup is needed for this and the database is growing, that can explain the non-linear course of the graph. This can be easily resolved by keeping the last hash value for currently connected peers in memory. However this is an indication that creating blocks by the thousands is an IO bound process, rather than CPU bound. Finally, if mobile devices are to be full-fletched nodes on the Tribler network, they should not slow down significantly more than any other ordinary laptop, besides being slower in the first place. Hardware acceleration could close this gap without sacrificing battery life too much. Because mobile devices are a bit behind on the technology curve with respect to desktop computers it is probable that the gap becomes smaller over the coming years. The capacity to store enough Multichain blocks to audit past exchanges should also be on par. If not, other more powerful nodes could be queried to supply the necessary history about a peer, that requests your bandwidth, to verify if that peer is trustworthy.

6.3. Startup time

Key to user retention is a fast startup time. Therefore the service starts up in the background, separately from the GUI. This way at the GUI is responsive to user input while the service may still be loading. One of the design objectives for separating the front-end from the back-end was responsiveness. However before any task can be executed by the service it needs to be actually started. To measure the startup time we register the time of launching the app and the moment the Tribler-started-event occurs. This event is sent over the API event-stream and signifies that Tribler is fully started and ready to accept all incoming requests. We expect consistent loading times because right after starting we shut Tribler down again. To get a good idea of how the user experience may differ we measure the startup time 10 times on 5 different devices. The app is launched with Android Debug Bridge (adb) from a laptop and the Tribler-started-event is read from logcat over adb and timed on the same laptop so they use the same clock.

| Device	N	Avg. (s)	Min. (s)	Max. (s)	s (s)
Nexus 10	10	3.781	3.416	4.085	0.211
Nexus 6	10	4.319	4.124	4.670	0.179
Nexus 5	10	3.353	3.273	3.459	0.081
Galaxy Nexus	10	7.086	6.161	7.772	0.454
S3	10	31.935	30.616	33.940	1.116

Table 6.1: ???

Figure 6.6: Startup times per device

Table 6.1 shows the statistics per device. The results show a very small sample standard deviation and a very low startup time. The S3 is performing worse than may be expected judging from the results of the other devices. The reason for that may be that this phone was not wiped and given a fresh install of Android. That could mean that other applications installed on a device could significantly impact the performance of Tribler. The sample standard deviation is relatively small though, which contradicts this hypothesis. This should be investigated further, including if anything can be done on the part of Tribler.

6.4. API latency

The user expects operations to take a consistent and reasonable amount of time. By design all functionality is going through the API. We use Apache JMeter to verify if the API is responding consistently within a reasonable amount of time by measuring the latency. JMeter measures the latency from just before sending the request to just after the first response has been received. [**?**] Thus the measurement includes the time needed to assemble the request as well as processing it and returning a response with latency on the way back. It excludes the transfer time of the complete response and subsequent processing and rendering time. Any client can process and render a response differently, possibly in a streaming fashion. By leaving out that element this metric is measuring the operation time via the API. We want to see that the latency is bounded, consistent and generally low. Previous research has shown the following results on a computer:

Figure 6.7: API latency on PC

We did the same benchmark with a slightly different setup. A Nexus 6 smart-phone with Android 6.0.1 Cyanogen mod was connected via USB to a laptop running JMeter. The API tcp-port was forwarded with ADB. With JMeter we request the discovered channels from the API a 1000 times and measure the latency. JMeter is also capable of setting a desired number of requests per minute, provided the device can handle it. We set this parameter to 300 requests per minute, equal to the other benchmark. Our database to be queried is not equal to one used in figure 6.7. The following figure shows the latency for every request.

	N	Avg. (ms)	Min. (ms)	Max. (ms)	σ (ms)	Throughput	KB/second	Avg. Bytes
Nexus 6	1000	1803	1292	3292	301.25	33.3 req/minute	347.44	641709.0
PC	1000	5	3	107	4.42	140.8 req/second	2548.00	18525.0

Table 6.2: These numbers are based on elapsed time, not latency, from the entire benchmark


Figure 6.8: API latency on Nexus 6 smart-phone

The plot shows large spikes and seemingly a curious and unexpected pattern. Table 6.2 shows the large difference in statistics. PC is added for reference. The set rate of 300 requests per minute is clearly not reached by the smart-phone in our benchmark by a factor of 9. The difference in amount of response data turned out to be much bigger than expected with a factor of 35. Tribler uses the event-driven networking engine Twisted, which is also written in Python. Twisted allows you to build inter-process communication protocols and provides a HTTP server which was used to build the REST API. Twisted uses a single thread to coordinate all others, called the reactor thread. If this thread is busy, the REST API can not receive incoming requests resulting in timeouts.

Considering the difference in amount of response data and the bandwidth of USB versus internal memory it is an unfair comparison. However, based on our result we may conclude that the API does not respond consistently and appear to behave differently than the other benchmark. The device not being able to process 5 requests per second possibly reveals information about the connection bandwidth and the multi-threaded processing capability. We assume that the 480MB/s theoretical bandwidth of USB2.0 is not a bottleneck considering the 2,5 MB/s result of the PC. Being a mobile device also other aspects my be at play here, like CPU frequency scaling. However this was turned off by acquiring a wake lock from the Android OS. It appears then that our design approach still suffers from performance issues due to imperfect multi-threaded performance. Because if the CPU is simply not powerful enough we expect to see a linear pattern instead of what we actually see in 6.8. This phenomenon is not of great importance though, because users are not expected to fire hundreds of requests per minute to the API. Further investigation should figure out if this phenomenon is also observed with lower amounts of requests per minute.

6.5. Profiling

Because of the challenges put forward in chapter 3.6 we investigate if time is spent disproportionately on some function. We expect that the limited resources of a mobile device may impact particular features more than others. If hardware acceleration is not present the less powerful CPU may struggle with encryption tasks. Instead of CPU time, time actually spent processing by the CPU, we measure wall-clock time. This way we measure the amount of time a user would have to wait for a certain function to be executed. We focus on wall clock time instead of CPU time because...

With the cProfile Python module and the visualisation tool SnakeViz we can see if any function takes a disproportionate amount of time. A Nexus 6 smart-phone with Android 6.0.1 Cyanogen mod was used for profiling Tribler. The profiler was running for 10 minutes with Tribler during normal operation and without any user input.


Figure 6.9: Profiler results of 10 minutes Tribler run time (bright pink represents business logic upon receiving a torrent)

Figure 6.9 shows that 27% of the time is spent on verifying cryptographic signatures. The bright pink represents the update function, which signifies various business logic upon receiving a torrent. Also notable is the same stack of functions within and outside of a community on top of the wrapper. This can be explained by the fact that torrents can be discovered outside of a community too. execute of apsw.Cursor execute of sqlite3.Cursor select.epoll [?] The time per call is most important for optimization because...

The significant chunk of time that the crypto takes is as expected. Since this task is actually delegated to the C library M2Crypto it should be possible to release the GIL of the Python interpreter so other Python code that does not depend on it can be executed. The main alternative provided in the standard library for CPU bound applications is the multiprocessing module, which works well for workloads that consist of relatively small numbers of long running computational tasks, but results in excessive message passing overhead if the duration of individual operations is short [?]. As seen from the time per call for __m2crypto.ecdsa_verify in table 6.3 the multiprocessing module would likely cause too much overhead. Another way to optimize is multiprocessing, which avoids the GIL completely [?].

# Calls	Total time (s)	Time per call (s)	Function
15	0.4867	0.03245	method 'commit' of 'sqlite3.Connection' objects
1	0.01692	0.01692	method 'executescript' of 'sqlite3.Cursor' objects
3820	64.28	0.01683	method 'poll' of 'select.epoll' objects
2	0.01048	0.005241	__m2crypto.ec_key_gen_key
31075	162	0.005212	__m2crypto.ecdsa_verify
1650	7.133	0.004323	__m2crypto.ecdsa_sign
1	0.001708	0.001708	_socket.gethostbyaddr
1	0.001284	0.001284	built-in method SSL_library_init
567	0.565	0.0009965	method 'executemany' of 'sqlite3.Cursor' objects
8	0.005731	0.0009552	__import__
12	0.01083	0.0009029	method 'connect_ex' of '_socket.socket' objects
1	0.00055	0.00055	built-in method SSL_load_error_strings
5	0.002515	0.000503	method 'recv' of '_socket.socket' objects
6989	3.436	0.0004917	method 'executemany' of 'apsw.Cursor' objects
1	0.000485	0.000485	dir
63677	20	0.000314	method 'execute' of 'apsw.Cursor' objects
5546	1.38	0.0002488	__m2crypto.ec_key_read_pubkey
15943	3.414	0.0002141	method 'sendto' of '_socket.socket' objects
44	0.009405	0.0002137	open
1	0.000212	0.000212	built-in method OpenSSL_add_all_algorithms
2	0.000405	0.0002025	__m2crypto.ec_key_new_by_curve_name
2	0.000394	0.000197	netifaces.interfaces
296	0.05179	0.000175	method 'sort' of 'list' objects
5	0.000826	0.0001652	__m2crypto.ec_key_read_bio
1	0.000147	0.000147	__m2crypto.rand_seed
4	0.00058	0.000145	netifaces.ifaddresses
47	0.005664	0.0001205	androidembed.log
12	0.001445	0.0001204	thread.start_new_thread
8	0.000936	0.000117	posix.mkdir
2240	0.2615	0.0001167	posix.open
17	0.001964	0.0001155	compile
3	0.00034	0.0001133	__m2crypto.ec_key_write_bio_no_cipher
5553	0.6196	0.0001116	__m2crypto.ec_key_write_pubkey
7	0.000777	0.000111	method 'send' of '_socket.socket' objects
140069	15.4	0.00011	method 'execute' of 'sqlite3.Cursor' objects
16	0.001759	0.0001099	method 'shutdown' of '_socket.socket' objects

Table 6.3: Native function calls wall clock time broken down per call during the 10 minute profiling (600 seconds total time)

6.6. CPU utilization

We saw in X that the cryptography takes a lot of time. We need to know if this is CPU bound or not for optimization purposes.

Python-for-Android supplies a CPython interpreter out of the box. CPython is optimized for single thread performance and compatibility with C extension modules. It is limited by a global interpreter lock (GIL) in multi-threaded use cases with shared memory. Tribler uses C extension modules for crypto tasks, which are CPU intensive. Tribler also uses the event-driven networking engine Twisted, which is written in Python. The core of the event loop within Twisted is the reactor, which runs on a single thread. The reactor provides a threading interface to offload long running tasks, such as IO or CPU intensive tasks to a thread pool. The GIL prohibits more than one thread to execute Python bytecode at a time. This negates all performance gains in terms of parallelism afforded by multi-core CPUs, making Python threads unusable for delegating CPU bound tasks to multiple cores. As shown in the precious section the crypto function took a considerable amount of time to compute. To see if the releasing the GIL as put forward as a solution is feasible we measure if the CPU has more capacity than is being utilized right now. During normal operation equivalent to the profiler measurement, we take a snapshot of the CPU utilization. If there is any performance to gain by releasing the GIL the CPU must be under-utilized right now. A Galaxy S3 smart-phone with Android 6.0.1 Cyanogen mod was used for this measurement.

Figure 6.10: CPU usage on a Galaxy S3 smart-phone, about 3 hours into normal operation

Figure 6.11: CPU usage on a Galaxy S3 smart-phone, about 2 minutes into the Multichain performance measurement

Run	Tests	Errors	Failures	Skipped	Line-rate
1	711	14	13	30	0.7241
2	749	12	15	3	0.7861

Table 6.4: Tests results and coverage line-rate at different points in time during development

The results show that indeed not al 4 cores of the CPU are utilized by a large margin. Even when performing the intensive crypto work of the Multichain experiment 2 CPU cores appear to be idling. This suggests that releasing the GIL during heavy crypto work could result in a significant performance gain. However our results are inconclusive and warrant further research.

6.7. HD video streaming

6.8. Testing and coverage

The design choice of reusing all Tribler core source code means we need to verify its correctness. To make sure all code on Android works the same as on other supported platforms we need to test all code. Tribler has some unit tests and integration tests that cover a large portion of the code, but not all. The ratio of tested lines of code with respect to the total number of lines of code is the coverage line-rate. We expect to see a line-rate value close to 1, but not 1 since we know the tests do not cover everything. All tests were run two times on the same device with 11 weeks of development in between. The nose module was used for running the tests together with the coverage module for gathering coverage data. The same Nexus 6 smart-phone with Android 6.0.1 Cyanogen mod was used in both runs. The following table shows the results of both executions. The number of tests has increased as well as the coverage line-rate while the number of errors and skipped tests have decreased. A failure means an assertion was not met and an error represents an exception while running a test. Therefore seeing that the number of failures increased is not that bad since the number of errors decreased. If the line-rate is 1 you still need the the branch-rate to be 1 as well before you can be confident the code will work as expected. The branch-rate is the number of code paths tested with respect to the total number of code paths possible. Unfortunately this metric is not part of the current test plan. Nevertheless the metrics show improvement overall.

7

Conclusions and future work

After validating and verifying our design in the previous chapter, we now come to the conclusions with regard to the main research question.

7.1. How feasible is it to run all Tribler functionality on mobile devices?

7.1.1. Given the constraints of mobile devices, what unique ability can be utilized to extend or enhance Tribler?

7.1.2. Using the unique properties of mobile devices, what features can be added or enhanced?

The feasibility in terms of functionality, performance and scalability were measured. Several of these measurements give raise to some minor considerations. The profiling revealed that crypto tasks are a significant part of processing messages. These tasks should be offloaded to a separate computational core to release the global interpreter lock. That would enable Python code to run in parallel, which in turn would improve performance and responsiveness. The CPU utilization measurement showed this approach is likely to succeed. The Multichain measurement showed that database performance in that instance was stringent in the current implementation. It is also the most easy to resolve by keeping the hash of the last block for connected peers in memory . All existing tests can be run on Android just as well as on other platforms. This means, in combination with the modularity of the new design and re-use of the Tribler core, that maintainability and testability are very fit. In terms of user experience the measurement of startup time showed to be very consistent and reasonable. The latency of the API however seemingly showed a curious pattern that cannot be explained right away. It is only known to occur if hundreds of requests per minute are fired at the API, which is unexpected behavior from the user.

Finally, to answer the secondary question: using the unique properties of mobile devices, what features can be added or enhanced? We added the capability to transfer the app to another device that doesn't have it yet via Bluetooth. If NFC is enabled on both devices the app can be transfered by just holding the devices next to each other. Also adding a channel to your favorites can be done this way. Using your real life social network you can build your on-line trusted network this way. And thanks to built-in hardware capability you can setup an ad hoc WiFi network to avoid any infrastructure completely.

What this means in the context of privacy and censorship is that if content is detected by the censor it may have already crossed or will cross the freedom border thanks to the properties of viral spreading.

7.2. Future work

This work enables a new direction for future research with Tribler: mobile devices. A new user group, that does not own desktop computers, gains access to Tribler's privacy enhancing functionality.

The contributions of this work enable future research into Tribler's intended use-cases fully geared towards mobile devices. Like for example a large scale experiment with various degrees of a powerful censor or a live deployment of Tribler mobile in areas with restricted Internet access to evaluate the effectiveness.

No routing, no blutooth p2p, mesh networking: that is done by Serval. Out of scope of this work. Add as enhancement later. We won't modify the networking capabilities based on the opportunities because that

would warrant a dedicated study.

In API version 18 support is added for Bluetooth Low Energy (LE), WiFi scan only mode, key store for app-private keys, hardware credential storage and automated UI testing.

The portability of our design and implementation allows it to be used in other applications like smart-tv's as well.

shared keychain, credit minging laptop, server, mobile..

NFC is used to share a channel id and to automatically setup a Bluetooth file transfer. In the future the same method can be used to exchange bootstrap peers for the Tribler network. Or integrate into apk just before sending via Bluetooth. The peer-to-peer Bluetooth file transfer could be made much faster if the WiFi Direct would be used.

Thumbnails retrieval and distribution now that video is so important on mobile, could be added on top of groundwork by me.

credit mining on device A, usage on device B

comparative analysis

Automate ad hoc WiFi direct like nfc-bluetooth!

cutting long running methods in smaller pieces, so multi-threading is smoother (only discovered now, earlier: why is gui hanging?) release GIL in native C code, mainly for heavy crypto tasks streaming API for big responses (dumping all known channels from db)

knippen voor twisted is gewenst, nu pas zichtbaar weakness in Tribler, nuttig om te verbeteren. (conclusion)

iOS with QtPython

Software development / technical aspect only Not policy making, organisational perspective, decision making, normative, ethical, Time limit of 9 months

refereren aan eigen werk in verbergen app en mensen beschermen op bepaalde manier combine with selfcompileapp

Bibliography

[1] S.M.A. Abbas, J.A. Pouwelse, D.H.J. Epema, and H.J. Sips. A gossip-based distributed social networking system. In *Enabling Technologies: Infrastructures for Collaborative Enterprises, 2009. WETICE '09. 18th IEEE International Workshops on*, pages 93–98, June 2009. doi: 10.1109/WETICE.2009.30. URL `http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=5159221&tag=1`.

[2] Erika Chin, Adrienne Porter Felt, Kate Greenwood, and David Wagner. Analyzing inter-application communication in android. In *Proceedings of the 9th international conference on Mobile systems, applications, and services - MobiSys '11*. Association for Computing Machinery (ACM), 2011. doi: 10.1145/1999995.2000018. URL `http://dx.doi.org/10.1145/1999995.2000018`.

[3] Roger Dingledine and Nick Mathewson. Design of a blocking-resistant anonymity system. *The Tor Project, Tech. Rep*, 1, 2006. URL `http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.101.7565&rep=rep1&type=pdf`.

[4] Roger Dingledine, Nick Mathewson, and Paul Syverson. Tor: The second-generation onion router. 2004. URL `http://www.dtic.mil/cgi-bin/GetTRDoc?Location=U2&doc=GetTRDoc.pdf&AD=ADA465464`.

[5] Gartner. Global mobile os market share in sales to end users from 1st quarter 2009 to 1st quarter 2016, 2016. URL `https://www.statista.com/statistics/266136/global-market-share-held-by-smartphone-operating-systems/`.

[6] Google. *C++ Library Support*, 2016. URL `https://developer.android.com/ndk/guides/cpp-support.html#sr`.

[7] Google. Platform versions, 2016. URL `https://developer.android.com/about/dashboards/index.html`.

[8] Google. *Intents and Intent Filters*, 2016. URL `https://developer.android.com/guide/components/intents-filters.html`.

[9] Dr David A. L. Levy, Nic Newman, Dr Richard Fletcher, and Dr Rasmus Kleis Nielsen. Digital news report 2016, 2016. URL `http://reutersinstitute.politics.ox.ac.uk/sites/default/files/Digital-News-Report-2016.pdf`.

[10] J.A. Pouwelse, P. Garbacki, J. Wang, A. Bakker, J. Yang, A. Iosup, D.H.J. Epema, M. Reinders, M. van Steen, and H.J. Sips. Tribler: A social-based peer-to-peer system. *Concurrency and Computation: Practice and Experience*, 20:127–138, February 2008. ISSN 1532-0634. URL `http://www.pds.ewi.tudelft.nl/pubs/papers/cpe2007.pdf`.

[11] Johan A. Pouwelse. The shadow internet: liberation from surveillance, censorship and servers. 2014. URL `https://tools.ietf.org/html/draft-pouwelse-perpass-shadow-internet-00`.

[12] R. Rahman, D. Hales, M. Meulpolder, V. Heinink, J. Pouwelse, and H. Sips. Robust vote sampling in a p2p media distribution system. In *Proceedings IPDPS 2009 (HotP2P 2009)*. IEEE Computer Society, May 2009. ISBN 978-1-4244-3750-4. URL `http://dx.doi.org/10.1109/IPDPS.2009.5160946`.

[13] Wendo Sabée, Dirk Schut, and Niels Spruit. Decentralized media streaming on android using tribler. 2014.

[14] Statista. Smartphone sales worldwide 2007-2015, 2016. URL `https://www.statista.com/statistics/263437/global-smartphone-sales-to-end-users-since-2007/`.

[15] C. van Bruggen, N. Feddes, and M. Vermeer. Anonymous hd video streaming for android using tribler, 2015.

[16] VideoLAN. Vlc for android, 2016. URL `https://www.videolan.org/vlc/download-android.html`.

[17] M.A. De Vos, R.M. Jagerman, and L.F.D. Versluis. Android tor tribler tunneling (at3). 2014. URL `http://repository.tudelft.nl/view/ir/uuid%3Ab258a5e8-002c-4631-9291-04be902119f6/`.